

Problem Set 4

AUTHOR
Sitong Guo, Helen Liu

PUBLISHED
November 3, 2024

PS4: Due Sat Nov 2 at 5:00PM Central. Worth 100 points. We use (*) to indicate a problem that we think might be time consuming.

Style Points (10 pts)

Please refer to the minilesson on code style [here](#).

Submission Steps (10 pts)

1. This problem set is a paired problem set.
2. Play paper, scissors, rock to determine who goes first. Call that person *Partner 1*.
 - Partner 1 (name and cnet ID): Sitong Guo (rehinkerg)
 - Partner 2 (name and cnet ID): Hailun Liu (hailunl)
3. Partner 1 will accept the `ps4` and then share the link it creates with their partner. You can only share it with one partner so you will not be able to change it after your partner has accepted.
4. "This submission is our work alone and complies with the 30538 integrity policy." Add your initials to indicate your agreement: `** SG ** ** HL **`
5. "I have uploaded the names of anyone else other than my partner and I worked with on the problem set [here](#)" (1 point)
6. Late coins used this pset: `** __ **` Late coins left after submission: `** __ **`
7. Knit your `ps4.qmd` to an PDF file to make `ps4.pdf`,
 - The PDF should not be more than 25 pages. Use `head()` and re-size figures when appropriate.
8. (Partner 1): push `ps4.qmd` and `ps4.pdf` to your github repo.
9. (Partner 1): submit `ps4.pdf` via Gradescope. Add your partner on Gradescope.
10. (Partner 1): tag your submission in Gradescope

Important: Repositories are for tracking code. **Do not commit the data or shapefiles to your repo.** The best way to do this is with `.gitignore`, which we have covered in class. If you do accidentally commit the data, Github has a [guide](#). The best course of action depends on whether you have pushed yet. This also means that both partners will have to download the initial raw data and any data cleaning code will need to be re-run on both partners' computers.

Download and explore the Provider of Services (POS) file (10 pts)

1. PRVDR_CTGRY_SBTYP_CD(Provider Category Subtype Code), PRVDR_CTGRY_CD(Provider Category Code), FAC_NAME(Facility Name), PRVDR_NUM(CMS Certification Number), PGM_TRMNTN_CD(Termination Code), ZIP_CD(Address: ZIP Code).
2. a. The file documented 7245 short term hospitals in 2016. This is likely to be overestimated for the acting quantity since that there is fraction of non-operating ones included.

```
::: {#6837d743 .cell execution_count=1} ``` {python .cell-code} import pandas as pd

pos2016 = pd.read_csv('E:/pos2016.csv', encoding='ISO-8859-1', dtype={'ZIP_CD': str})

short_term_hospitals = pos2016[(pos2016['PRVDR_CTGRY_CD'] == 1) &
(pos2016['PRVDR_CTGRY_SBTYP_CD'] == 1)]

num_hospitals = short_term_hospitals['PRVDR_NUM'].nunique()

print(f"Number of short-term hospitals in 2016: {num_hospitals}") ```

::: {cell-output .cell-output-stderr} ``` C:\18256\3682174362.py:3: DtypeWarning:

Columns (3) have mixed types. Specify dtype option on import or set low_memory=False.

:::

::: {cell-output .cell-output-stdout}
```

Number of short-term hospitals in 2016: 7245 `` `` `` ``

- b. According to the KFF, as of Jul 07, 2016, there are nearly 5,000 short-term, acute care hospitals in the United States. As stated by 2016 CMS Statistics, on the other hand, Medicare short-term hospital was only 3436. This discrepancy is due to that the file contains hospitals not acting and there might be other short-term hospitals besides acute care. (The Number of U.S. Hospitals by Type (Total 5534), FY2016 by American Hospital Association claimed that there were only 5534 hospitals in total in 2016, this is due to some inexplicable divergences on statistical caliber on definition of hospital.)

3.

```
import matplotlib.pyplot as plt
import altair as alt

pos2017 = pd.read_csv('E:/pos2017.csv', encoding='ISO-8859-1', dtype={'ZIP_CD': str})
pos2018 = pd.read_csv('E:/pos2018.csv', encoding='ISO-8859-1', dtype={'ZIP_CD': str})
pos2019 = pd.read_csv('E:/pos2019.csv', encoding='ISO-8859-1', dtype={'ZIP_CD': str})

def short_term(df):
    return df[(df['PRVDR_CTGRY_CD'] == 1) & (df['PRVDR_CTGRY_SBTYP_CD'] == 1)]

short_term_2016 = short_term(pos2016)
short_term_2017 = short_term(pos2017)
short_term_2018 = short_term(pos2018)
short_term_2019 = short_term(pos2019)

short_term_2016['Year'] = 2016
short_term_2017['Year'] = 2017
short_term_2018['Year'] = 2018
short_term_2019['Year'] = 2019

# append
all_years = pd.concat([short_term_2016, short_term_2017, short_term_2018, short_term_2019])

counts_by_year = all_years['Year'].value_counts().sort_index().reset_index()
counts_by_year.columns = ['Year', 'Count']

# set range for the y axis
y_min = counts_by_year['Count'].min() - 10
y_max = counts_by_year['Count'].max() + 10

alt.Chart(counts_by_year).mark_line(color='lightblue').encode(
    x=alt.X('Year:O', title='Year'),
    y=alt.Y('Count:Q', title='Number of Short-Term Hospitals', scale=alt.Scale(domain=[y_min, y_max]))
).properties(
    title='Number of Short-Term Hospital Observations by Year',
    height = 400,
    width = 400
)
```

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\3356294368.py:16: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\3356294368.py:17: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\3356294368.py:18: SettingWithCopyWarning:

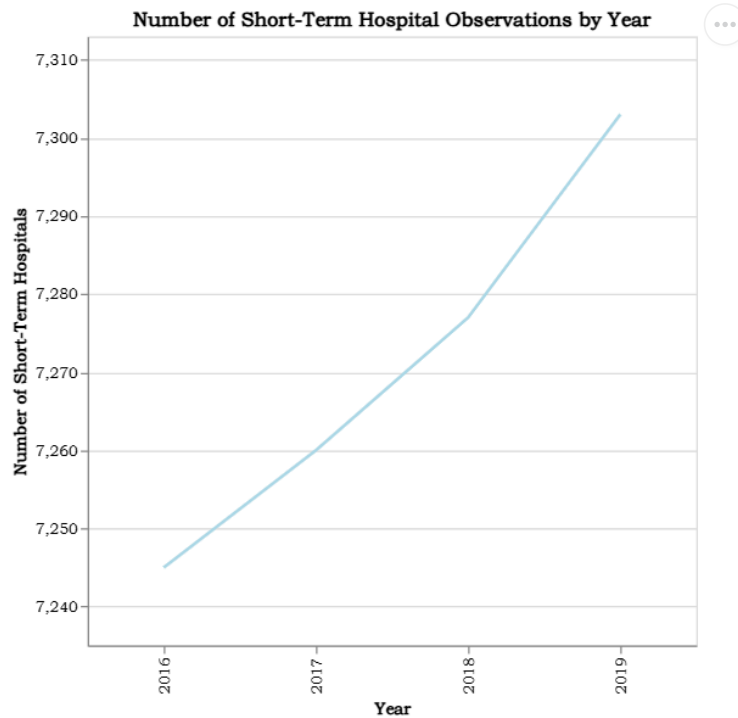
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\3356294368.py:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy



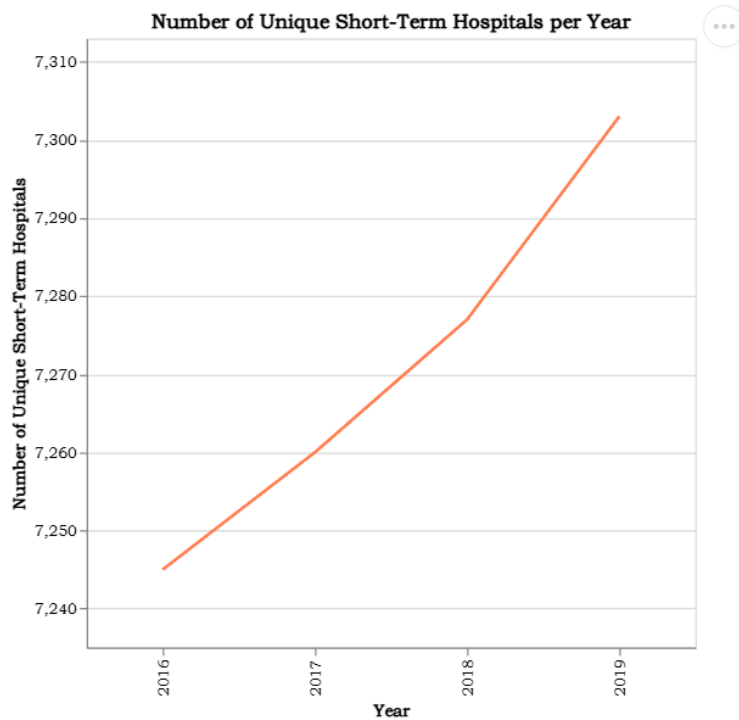
4. a.

```
::: {#a2ba149c .cell execution_count=3} ``` {python .cell-code} unique_hospitals_per_year =  
all_years.groupby("Year")["PRVDR_NUM"].nunique().reset_index() unique_hospitals_per_year.columns =  
['Year', 'Unique_Count']
```

```
y_min = unique_hospitals_per_year['Unique_Count'].min() - 10 y_max =  
unique_hospitals_per_year['Unique_Count'].max() + 10
```

```
alt.Chart(unique_hospitals_per_year).mark_line(color='coral').encode( x=alt.X('Year:O', title='Year'),  
y=alt.Y('Unique_Count:Q', title='Number of Unique Short-Term Hospitals', scale=alt.Scale(domain=  
[y_min, y_max])) ).properties( title = 'Number of Unique Short-Term Hospitals per Year', width = 400,  
height = 400 ) ```
```

```
::: {cell-output .cell-output-display execution_count=3}
```



... ..

b. They are identical, which means that the short-term hospitals were each having a unique CCN without redundancy, in the period of 2016-2019.

Identify hospital closures in POS file (15 pts) (*)

1.

```
pos20161 = pos2016[pos2016['PRVDR_CTGRY_CD'] == 1]
pos20161 = pos20161[pos20161['PRVDR_CTGRY_SBTYP_CD'] == 1]
pos20171=pos2017[pos2017['PRVDR_CTGRY_CD'] == 1]
pos20171=pos20171[pos20171['PRVDR_CTGRY_SBTYP_CD'] == 1]
pos20181=pos2018[pos2018['PRVDR_CTGRY_CD'] == 1]
pos20181=pos20181[pos20181['PRVDR_CTGRY_SBTYP_CD'] == 1]
pos20191=pos2019[pos2019['PRVDR_CTGRY_CD'] == 1]
pos20191=pos20191[pos20191['PRVDR_CTGRY_SBTYP_CD'] == 1]

pos2016_subset = pos20161[['FAC_NAME', 'PRVDR_NUM', 'PGM_TRMNTN_CD', 'ZIP_CD']]
pos2017_subset = pos20171[['FAC_NAME', 'PRVDR_NUM', 'PGM_TRMNTN_CD', 'ZIP_CD']]
pos2018_subset = pos20181[['FAC_NAME', 'PRVDR_NUM', 'PGM_TRMNTN_CD', 'ZIP_CD']]
pos2019_subset = pos20191[['FAC_NAME', 'PRVDR_NUM', 'PGM_TRMNTN_CD', 'ZIP_CD']]

pos2016_subset = pos2016_subset.add_suffix('_2016').rename(columns={'PRVDR_NUM_2016': 'PRVDR_NUM'})
pos2017_subset = pos2017_subset.add_suffix('_2017').rename(columns={'PRVDR_NUM_2017': 'PRVDR_NUM'})
pos2018_subset = pos2018_subset.add_suffix('_2018').rename(columns={'PRVDR_NUM_2018': 'PRVDR_NUM'})
pos2019_subset = pos2019_subset.add_suffix('_2019').rename(columns={'PRVDR_NUM_2019': 'PRVDR_NUM'})

merged = pos2016_subset.merge(pos2017_subset, on='PRVDR_NUM', how='outer') \
    .merge(pos2018_subset, on='PRVDR_NUM', how='outer') \
    .merge(pos2019_subset, on='PRVDR_NUM', how='outer')

merged_2016_0 = merged[merged['PGM_TRMNTN_CD_2016'] == 0]

suspected = merged_2016_0[
    ~((merged_2016_0['PGM_TRMNTN_CD_2016'] == 0) &
      (merged_2016_0['PGM_TRMNTN_CD_2017'] == 0) &
      (merged_2016_0['PGM_TRMNTN_CD_2018'] == 0) &
      (merged_2016_0['PGM_TRMNTN_CD_2019'] == 0))
]
suspected = suspected[['PRVDR_NUM', 'FAC_NAME_2016', 'ZIP_CD_2016', 'PGM_TRMNTN_CD_2016', 'PGM_TRMNTN_CD_2017', 'PGM_TRMNTN_CD_2018', 'PGM_TRMNTN_CD_2019']]

count = suspected['FAC_NAME_2016'].unique()
print(len(count))
```

There are 174 hospital that fit this definition(hospitals that were active in 2016 that were suspected to have closed by 2019)

2.

```
import numpy as np
conditions = [
    (suspected['PGM_TRMNTN_CD_2017'] != 0) & (suspected['PGM_TRMNTN_CD_2018'] != 0) & (suspected['PGM_TRMNTN_CD_2019'] != 0) &
    (suspected['PGM_TRMNTN_CD_2017'] == 0) & (suspected['PGM_TRMNTN_CD_2018'] != 0) & (suspected['PGM_TRMNTN_CD_2019'] != 0) &
    (suspected['PGM_TRMNTN_CD_2017'] == 0) & (suspected['PGM_TRMNTN_CD_2018'] == 0) & (suspected['PGM_TRMNTN_CD_2019'] != 0)
]
choices = [2017, 2018, 2019]
suspected['YearOfSuspectedClosure'] = np.select(conditions, choices, default=np.nan)

sorted_suspected = suspected.sort_values(by='FAC_NAME_2016')[['FAC_NAME_2016', 'YearOfSuspectedClosure',
                                                                'PGM_TRMNTN_CD_2016', 'PGM_TRMNTN_CD_2017', 'PGM_TRMNTN_CD_2018', 'PGM_TRMNTN_CD_2019']]
top_10_hospitals = sorted_suspected.head(10)
print(top_10_hospitals)
```

	FAC_NAME_2016	YearOfSuspectedClosure
168	ABRAZO MARYVALE CAMPUS	2017.0
568	ADVENTIST MEDICAL CENTER - CENTRAL VALLEY	2017.0
4852	AFFINITY MEDICAL CENTER	2018.0
4356	ALBANY MEDICAL CENTER / SOUTH CLINICAL CAMPUS	2017.0
6273	ALLEGIANCE SPECIALTY HOSPITAL OF KILGORE	2017.0
3596	ALLIANCE LAIRD HOSPITAL	2019.0
4990	ALLIANCEHEALTH DEACONESS	2019.0
1384	ANNE BATES LEACH EYE HOSPITAL	2019.0
1044	ARKANSAS VALLEY REGIONAL MEDICAL CENTER	2017.0
3975	BANNER CHURCHILL COMMUNITY HOSPITAL	2017.0

	ZIP_CD_2016
168	85031
568	93230
4852	44646
4356	12208
6273	75662
3596	39365
4990	73112
1384	33136
1044	81050
3975	89406

3. a.

```
::: {#92bb617b .cell execution_count=6} ``` {python .cell-code} #hospitals active each year
active_2016 = pos2016_subset[pos2016_subset['PGM_TRMNTN_CD_2016'] == 0][['ZIP_CD_2016', 'PRVDR_NUM']]
active_2017 = pos2017_subset[pos2017_subset['PGM_TRMNTN_CD_2017'] == 0][['ZIP_CD_2017', 'PRVDR_NUM']]
active_2018 = pos2018_subset[pos2018_subset['PGM_TRMNTN_CD_2018'] == 0][['ZIP_CD_2018', 'PRVDR_NUM']]
active_2019 = pos2019_subset[pos2019_subset['PGM_TRMNTN_CD_2019'] == 0][['ZIP_CD_2019', 'PRVDR_NUM']]
```

```
#by ZIP each year active_count_2016 =
active_2016.groupby('ZIP_CD_2016').size().reset_index(name='ActiveCount_2016')
active_count_2017 = active_2017.groupby('ZIP_CD_2017').size().reset_index(name='ActiveCount_2017')
active_count_2018 = active_2018.groupby('ZIP_CD_2018').size().reset_index(name='ActiveCount_2018')
active_count_2019 = active_2019.groupby('ZIP_CD_2019').size().reset_index(name='ActiveCount_2019')
```

```
active_count_2016 = active_count_2016.rename(columns={'ZIP_CD_2016': 'ZIP_CD'})
active_count_2017 = active_count_2017.rename(columns={'ZIP_CD_2017': 'ZIP_CD'})
active_count_2018 = active_count_2018.rename(columns={'ZIP_CD_2018': 'ZIP_CD'})
active_count_2019 = active_count_2019.rename(columns={'ZIP_CD_2019': 'ZIP_CD'})
```

```
# merge active counts to compare by ZIP
zip_counts = active_count_2016.merge(active_count_2017, on='ZIP_CD', how='outer')
zip_counts.merge(active_count_2018, on='ZIP_CD', how='outer')
zip_counts.merge(active_count_2019, on='ZIP_CD', how='outer')
```

```
suspected = suspected.merge(zip_counts, left_on='ZIP_CD_2016', right_on='ZIP_CD', how='left')
```

```
#number of active did not decrease...questionable! suspected['IsMerger'] = np.where(
((suspected['YearOfSuspectedClosure'] == 2017) & (suspected['ActiveCount_2017'] >=
suspected['ActiveCount_2016'])) | ((suspected['YearOfSuspectedClosure'] == 2018) &
(suspected['ActiveCount_2018'] >= suspected['ActiveCount_2017'])) |
((suspected['YearOfSuspectedClosure'] == 2019) & (suspected['ActiveCount_2019'] >=
suspected['ActiveCount_2018'])), True, False)
```

```
#potential merger/acquisition corrected_closures = suspected[~suspected['IsMerger']] merger_count =
suspected['IsMerger'].sum() print(f"potential mergers/acquisitions: {merger_count}") ``
```

```
::: {cell-output .cell-output-stdout} potential mergers/acquisitions: 6 ::: :::
```

b.

```
::: {#bbe9c8d7 .cell execution_count=7} {.python .cell-code} #corrected closures
corrected_closure_count = corrected_closures['FAC_NAME_2016'].nunique() print(f"hospitals correcting
for m/a: {corrected_closure_count}")
```

```
::: {cell-output .cell-output-stdout} hospitals correcting for m/a: 168 ::: :::
```

c.

```
::: {#145ce8fe .cell execution_count=8} {.python .cell-code} sorted_corrected_closures =
corrected_closures.sort_values(by='FAC_NAME_2016')[['FAC_NAME_2016', 'YearOfSuspectedClosure',
'ZIP_CD_2016']] print(sorted_corrected_closures.head(10))
```

```
::: {cell-output .cell-output-stdout} `` FAC_NAME_2016 YearOfSuspectedClosure
4 ABRAZO MARYVALE CAMPUS 2017.0
10 ADVENTIST MEDICAL CENTER - CENTRAL VALLEY 2017.0
97 AFFINITY MEDICAL CENTER 2018.0
80 ALBANY MEDICAL CENTER / SOUTH CLINICAL CAMPUS 2017.0
140 ALLEGIANCE SPECIALTY HOSPITAL OF KILGORE 2017.0
62 ALLIANCE LAIRD HOSPITAL 2019.0
101 ALLIANCEHEALTH DEACONESS 2019.0
26 ANNE BATES LEACH EYE HOSPITAL 2019.0
21 ARKANSAS VALLEY REGIONAL MEDICAL CENTER 2017.0
69 BANNER CHURCHILL COMMUNITY HOSPITAL 2017.0
```

```
ZIP_CD_2016
```

```
4 85031
10 93230
97 44646
80 12208
140 75662
62 39365
101 73112
26 33136
21 81050
69 89406
`` ::: :::
```

Download Census zip code shapefile (10 pt)

1. a. (1).shp (Shape file): Main file that has feature geometrics, such as points, lines, or polygons that represent the shapes of geographic features including ZIP code boundaries.
- (2).shx (Shape index file) : Contains an positional index of the geometries in the shp file, accelerating access to geographic features.
- (3).dbf (database file): a tabular file with attribute information, in dBASE format that stores attributes or additional data about each shape in the shp file.
- (4).prj (projection file): Describes the Coordinate Reference System (CRS). Contains information about the system and projection used in the shp, ensuring that the data aligns correctly with other geographic data.

(5).xml: Detailed text information about the dataset like source, description, date, attribute definitions, and other information for understanding the data's context and structure.

b.

```
::: {#eb81ee47 .cell execution_count=9} ```{.python .cell-code} import os
directory = 'E:/SeriousBusiness/Applications/uchicago/python2'
```

```
for filename in os.listdir(directory):
    file_path = os.path.join(directory, filename)
    if os.path.isfile(file_path):
        file_size = os.path.getsize(file_path)
        print(f"{filename}: {file_size} KB")
```

```
::: {cell-output .cell-output-stdout} gz_2010_us_860_00_500k.dbf: 6425474 KB
gz_2010_us_860_00_500k.prj: 165 KB gz_2010_us_860_00_500k.shp: 837544580 KB
gz_2010_us_860_00_500k.shx: 265060 KB gz_2010_us_860_00_500k.xml: 15639 KB PS4.docx: 39891 KB
PS4.pdf: 240979 KB pset4_template.qmd: 3234 KB ~$PS4.docx: 162 KB ::::
```

2.

```
import geopandas as gpd
import pandas as pd
import matplotlib.pyplot as plt
```

```
#number of hospitals per ZIP
hospitals_per_zip = active_2016.groupby('ZIP_CD_2016').size().reset_index(name='Hospital_Count')
hospitals_per_zip.columns = ['ZIP_CD', 'Hospital_Count']
```

```
zip_codes = gpd.read_file('E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.shp')
```

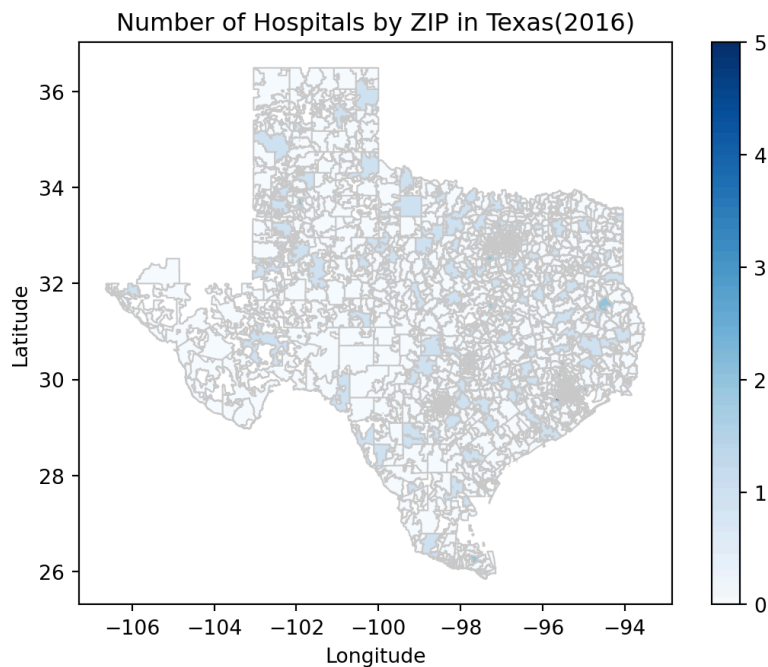
```
#print(zip_codes.columns)
# Index(['GEO_ID', 'ZCTA5', 'NAME', 'LSAD', 'CENSUSAREA', 'geometry'], dtype='object')
#Texas
texas_zip_codes = zip_codes[zip_codes['ZCTA5'].str.startswith(('75', '76', '77', '78', '79'))]
```

```
texas_zip_hospitals = texas_zip_codes.merge(hospitals_per_zip, left_on='ZCTA5', right_on='ZIP_CD', how='left')
#fill missing hospital counts with 0!
texas_zip_hospitals['Hospital_Count'] = texas_zip_hospitals['Hospital_Count'].fillna(0)
texas_zip_hospitals.head(5)
```

	GEO_ID	ZCTA5	NAME	LSAD	CENSUSAREA	geometry	ZIP_CD	Hospital_Count
0	8600000US78624	78624	78624	ZCTA5	708.041	POLYGON ((-98.96423 30.49848, -98.96416 30.498...	78624	1.0
1	8600000US78626	78626	78626	ZCTA5	93.046	POLYGON ((-97.60944 30.57185, -97.61688 30.568...	NaN	0.0
2	8600000US78628	78628	78628	ZCTA5	73.382	POLYGON ((-97.69285 30.57122, -97.69286 30.571...	NaN	0.0
3	8600000US78631	78631	78631	ZCTA5	325.074	POLYGON ((-99.13053 30.36555, -99.13065 30.365...	NaN	0.0
4	8600000US78632	78632	78632	ZCTA5	96.278	POLYGON ((-97.40946 29.75929, -97.40947 29.758...	NaN	0.0

```
plt.figure(figsize=(18, 15))
texas_zip_hospitals.plot(column='Hospital_Count', cmap='Blues', linewidth=0.8, edgecolor='0.8', legend=True)
plt.title('Number of Hospitals by ZIP in Texas(2016)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```

<Figure size 1728x1440 with 0 Axes>



Calculate zip code's distance to the nearest hospital (20 pts) (*)

1.

```
shp = gpd.read_file('E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.shp')
shx = gpd.read_file('E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.shx')
zips_all_centroids = shp.copy()
zips_all_centroids['geometry'] = shp.geometry.centroid
```

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\65589225.py:4: UserWarning:

Geometry is in a geographic CRS. Results from 'centroid' are likely incorrect. Use 'GeoSeries.to_crs()' to re-project geometries to a projected CRS before this operation.

Dimensions: The dimensions output from the GeoDataFrame created for the centroid of each zip code nationally gives the number of rows and columns which is equal to the number of unique ZIP codes in the data. Each row represents a unique ZIP Code area, and the columns represent the . The coordinate position of the centroids(each ZIP Code).

Columns meaning: 1.GEO_ID: A unique identifier for each geographic region, helping to distinguish each ZIP Code area. 2.ZCTA5:This is the 5-digit ZIP Code Tabulation Area (ZCTA), a representation of ZIP Code regions used in census data. 3.NAME: Typically represents the ZIP Code for the area. 4.LSAD: Legal/Statistical Area Description, identifying the type of area, such as city or rural. 5.geometry: Contains the centroid point of each ZIP Code area.This is a Point geometry representing the geometric center of each ZIP Code area.

2. In subsets of zips_texas_centroids,there are 1935 unique zip codes. In subsets of zips_texas_borderstates_centroids,there are 4057 unique zip codes.

```
import geopandas as gpd
zips_all_centroids['ZIP_INT'] = zips_all_centroids['ZCTA5'].astype(int)
texas_condition = (
    (zips_all_centroids['ZIP_INT'] >= 75000) & (zips_all_centroids['ZIP_INT'] <= 79999)
)

border_states_condition = (
    ((zips_all_centroids['ZIP_INT'] >= 87000) & (zips_all_centroids['ZIP_INT'] <= 88499)) |
    ((zips_all_centroids['ZIP_INT'] >= 70000) & (zips_all_centroids['ZIP_INT'] <= 72999)) |
    ((zips_all_centroids['ZIP_INT'] >= 73000) & (zips_all_centroids['ZIP_INT'] <= 79999))
)
zips_texas_centroids = zips_all_centroids[texas_condition]
zips_texas_borderstates_centroids = zips_all_centroids[texas_condition | border_states_condition]
```



```
unique_texas_zips = zips_texas_centroids['ZCTA5'].nunique()
unique_texas_border_zips = zips_texas_borderstates_centroids['ZCTA5'].nunique()
print("Unique ZIP Codes in Texas:", unique_texas_zips)
print("Unique ZIP Codes in Texas and bordering states:", unique_texas_border_zips)
```

Unique ZIP Codes in Texas: 1935

Unique ZIP Codes in Texas and bordering states: 4057

3. There are 468 rows in `zips_withhospital_centroids`.

We'll use an inner join since we only want the zip codes that appear in both datasets (those in `zips_texas_borderstates_centroids` that are also in `merged_gdf`). The merge will be based on the zip code column, typically named something like 'ZIP_CD_2016' in both GeoDataFrames. Ensure both columns are named consistently for the merge.

```
dbf = gpd.read_file('E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.dbf')
merged_2016_0['ZIP_INT'] = pd.to_numeric(merged_2016_0['ZIP_CD_2016'], errors='coerce')
border_states_condition = (
    ((merged_2016_0['ZIP_INT'] >= 87000) & (merged_2016_0['ZIP_INT'] <= 88499)) |
    ((merged_2016_0['ZIP_INT'] >= 70000) & (merged_2016_0['ZIP_INT'] <= 72999)) |
    ((merged_2016_0['ZIP_INT'] >= 73000) & (merged_2016_0['ZIP_INT'] <= 79999))
)
border_states_hospital_2016 = merged_2016_0[border_states_condition]
border_states_hospital_2016 = border_states_hospital_2016.drop(columns=['ZIP_INT'])

texas_borderstates_hospital = border_states_hospital_2016.groupby('ZIP_CD_2016').size().reset_index(name='count')
texas_borderstates_hospital['ZIP_CD_2016'] = texas_borderstates_hospital['ZIP_CD_2016'].astype(int)
texas_borderstates_hospital['ZIP_CD_2016'] = texas_borderstates_hospital['ZIP_CD_2016'].astype(str)

df = pd.DataFrame(zips_texas_borderstates_centroids.drop(columns='geometry'))
df = df.rename(columns={'ZCTA5': 'ZIP_CD_2016'})
df['ZIP_CD_2016'] = df['ZIP_CD_2016'].astype(str)
```

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\2794894124.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
8870    70003
8871    70030
8872    70032
8873    70036
8874    70038
...
32917   78261
32918   78368
32919   78412
32920   78557
32921   78586
```

Name: ZIP_CD_2016, Length: 4057, dtype: object

```
merged_gdf = df.merge(
    texas_borderstates_hospital,
    on='ZIP_CD_2016',
    how='right'
)
print(f"There are {len(merged_gdf)} rows in zips_withhospital_centroids .")
```

There are 468 rows in `zips_withhospital_centroids`.

```
zips_withhospital_centroids = gpd.GeoDataFrame(merged_gdf)
```

4. a. It will take about 0.0099 seconds for subset to 10 zip codes. And the whole process will take about 1.93 seconds.

```
import time
#shp = gpd.read_file('E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.shp')
zips_texas_centroids = shp[shp['ZCTA5'].astype(str).str[:2].isin(['75', '76', '77', '78', '79'])]
zips_withhospital_centroids = shp[shp['ZCTA5'].isin(merged_gdf['ZIP_CD_2016'])]

zips_texas_centroids['geometry'] = zips_texas_centroids.geometry.centroid
zips_withhospital_centroids['geometry'] = zips_withhospital_centroids.geometry.centroid

zips_texas_subset = zips_texas_centroids.head(10)
start_time = time.time()
subset_join_result = gpd.sjoin_nearest(
    zips_texas_subset,
    zips_withhospital_centroids,
    how="inner",
    distance_col="distance"
)
time_taken = time.time() - start_time
subset_join_result, time_taken
```

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\196238175.py:6: UserWarning:

Geometry is in a geographic CRS. Results from 'centroid' are likely incorrect. Use 'GeoSeries.to_crs()' to re-project geometries to a projected CRS before this operation.

E:\anaconda3\Lib\site-packages\geopandas\geodataframe.py:1819: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\RedthinkerDantler\AppData\Local\Temp\ipykernel_18256\196238175.py:7: UserWarning:

Geometry is in a geographic CRS. Results from 'centroid' are likely incorrect. Use 'GeoSeries.to_crs()' to re-project geometries to a projected CRS before this operation.

E:\anaconda3\Lib\site-packages\geopandas\array.py:403: UserWarning:

Geometry is in a geographic CRS. Results from 'sjoin_nearest' are likely incorrect. Use 'GeoSeries.to_crs()' to re-project geometries to a projected CRS before this operation.

(GEO_ID_left	ZCTA5_left	NAME_left	LSAD_left	CENSUSAREA_left	\
9207	8600000US78624	78624	78624	ZCTA5	708.041	
9208	8600000US78626	78626	78626	ZCTA5	93.046	
9209	8600000US78628	78628	78628	ZCTA5	73.382	
9210	8600000US78631	78631	78631	ZCTA5	325.074	
9211	8600000US78632	78632	78632	ZCTA5	96.278	
9212	8600000US78633	78633	78633	ZCTA5	82.269	
9213	8600000US78634	78634	78634	ZCTA5	63.656	
9214	8600000US78635	78635	78635	ZCTA5	15.940	
9215	8600000US78636	78636	78636	ZCTA5	349.689	
9216	8600000US78638	78638	78638	ZCTA5	114.562	

	geometry	index_right	GEO_ID_right	ZCTA5_right	\
9207	POINT (-98.87707 30.2816)	9207	8600000US78624	78624	
9208	POINT (-97.59733 30.66535)	24856	8600000US78664	78664	
9209	POINT (-97.75112 30.64108)	9231	8600000US78681	78681	
9210	POINT (-99.30528 30.33772)	32903	8600000US78028	78028	
9211	POINT (-97.47045 29.69633)	32565	8600000US78629	78629	
9212	POINT (-97.75426 30.74197)	9231	8600000US78681	78681	
9213	POINT (-97.54471 30.55908)	24856	8600000US78664	78664	
9214	POINT (-98.55961 30.21086)	9207	8600000US78624	78624	
9215	POINT (-98.41885 30.30504)	9223	8600000US78654	78654	
9216	POINT (-97.79495 29.6569)	26359	8600000US78155	78155	

	NAME_right	LSAD_right	CENSUSAREA_right	distance
9207	78624	ZCTA5	708.041	0.000000

9208	78664	ZCTA5	16.562	0.167651
9209	78681	ZCTA5	21.727	0.110844
9210	78028	ZCTA5	250.675	0.337251
9211	78629	ZCTA5	425.389	0.219909
9212	78681	ZCTA5	21.727	0.210633
9213	78664	ZCTA5	16.562	0.114657
9214	78624	ZCTA5	708.041	0.325246
9215	78654	ZCTA5	200.189	0.340902
9216	78155	ZCTA5	354.566	0.189651

0.0866093635559082)

b.As for the full calculation,it takes about 0.07 seconds,which is faster than 1 estimated.

```

zips_texas_centroids = zips_texas_centroids.to_crs(epsg=3857)
zips_withhospital_centroids = zips_withhospital_centroids.to_crs(epsg=3857)
start_time = time.time()
subset_join_result = gpd.sjoin_nearest(
    zips_texas_centroids,
    zips_withhospital_centroids,
    how="inner",
    distance_col="distance"
)
time_taken = time.time() - start_time
subset_join_result, time_taken

```

(GEO_ID_left	ZCTA5_left	NAME_left	LSAD_left	CENSUSAREA_left	\
9207	8600000US78624	78624	78624	ZCTA5	708.041	
9208	8600000US78626	78626	78626	ZCTA5	93.046	
9209	8600000US78628	78628	78628	ZCTA5	73.382	
9210	8600000US78631	78631	78631	ZCTA5	325.074	
9211	8600000US78632	78632	78632	ZCTA5	96.278	
...	
32917	8600000US78261	78261	78261	ZCTA5	29.865	
32918	8600000US78368	78368	78368	ZCTA5	216.341	
32919	8600000US78412	78412	78412	ZCTA5	8.798	
32920	8600000US78557	78557	78557	ZCTA5	11.653	
32921	8600000US78586	78586	78586	ZCTA5	176.313	

	geometry	index_right	GEO_ID_right	\
9207	POINT (-11006945.524 3539798.214)	9207	8600000US78624	
9208	POINT (-10864484.611 3589364.081)	24856	8600000US78664	
9209	POINT (-10881605.454 3586223.598)	9231	8600000US78681	
9210	POINT (-11054613.303 3547034.688)	32903	8600000US78028	
9211	POINT (-10850361.016 3464574.867)	32565	8600000US78629	
...	
32917	POINT (-10954048.141 3463994.843)	10809	8600000US78258	
32918	POINT (-10888193.07 3262273.761)	26302	8600000US78102	
32919	POINT (-10836170.333 3211684.17)	32919	8600000US78412	
32920	POINT (-10936401.874 3012307.913)	10839	8600000US78503	
32921	POINT (-10868244.776 3012099.803)	10847	8600000US78550	

	ZCTA5_right	NAME_right	LSAD_right	CENSUSAREA_right	distance
9207	78624	78624	ZCTA5	708.041	0.000000
9208	78664	78664	ZCTA5	16.562	21443.710461
9209	78681	78681	ZCTA5	21.727	14228.011802
9210	78028	78028	ZCTA5	250.675	42374.908087
9211	78629	78629	ZCTA5	425.389	28125.820177
...
32917	78258	78258	ZCTA5	15.874	12854.585280
32918	78102	78102	ZCTA5	469.003	39136.894333
32919	78412	78412	ZCTA5	8.798	0.000000
32920	78503	78503	ZCTA5	17.036	7258.416063
32921	78550	78550	ZCTA5	95.609	19219.631838

[1935 rows x 13 columns],
0.026610136032104492)

c.The .prj file specifies that the unit is in "Degree" (angular unit)

In this context, "Degree" represents degrees of latitude and longitude. On the Earth's surface, 1 degree of latitude is approximately equal to 69 miles. The distance represented by 1 degree of longitude varies by latitude, but in mid-latitude regions like the United States, 1 degree of longitude is roughly 53 miles. So, the

approximate conversions are: 1 degree of latitude \approx 69 miles 1 degree of longitude \approx 53 miles (suitable for mid-latitude areas in the U.S.)

```
from pyproj import CRS
prj_file = "E:/SeriousBusiness/Applications/uchicago/python2/gz_2010_us_860_00_500k.prj"
with open(prj_file, 'r') as file:
    prj_text = file.read()
crs = CRS.from_wkt(prj_text)
print(crs)
```

```
GEOGCS["GCS_North_American_1983",DATUM["D_North_American_1983",SPHEROID["GRS_1980",6378137,298.257222101
]],PRIMEM["Greenwich",0],UNIT["Degree",0.017453292519943295]]
```

5. a. Unit is included distance.

```
subset_join_result = gpd.sjoin_nearest(
    zips_texas_centroids,
    zips_withhospital_centroids,
    how="inner",
    distance_col="distance"
)
```

b. the average distance is 0.21101748566398393.

the average distance in miles is 14.49 miles


```
dd = subset_join_result.drop(columns='geometry')
mean_distance = dd['distance'].mean()
print(mean_distance)
```

25385.83933214372

c.

```
import matplotlib.pyplot as plt
subset_join_result.plot(column = 'distance').set_axis_off()
plt.axis("off")
```

```
(-11939710.352867601,
-10348254.934308909,
2924394.112609078,
4435184.103883993)
```



Effects of closures on access in Texas (15 pts)

1.

```

texas_closures = sorted_corrected_closures[sorted_corrected_closures['ZIP_CD_2016'].str.startswith(('75',
closures_by_zip = texas_closures.groupby('ZIP_CD_2016').size().reset_index(name='Number_of_Closures')
closures_by_zip.columns = ['ZIP Code', 'Number of Closures']
print(closures_by_zip)

```

	ZIP Code	Number of Closures
0	75042	1
1	75051	1
2	75087	1
3	75140	1
4	75231	1
5	75235	1
6	75390	1
7	75601	1
8	75662	1
9	75835	1
10	75862	1
11	76502	1
12	76520	1
13	76531	1
14	76645	1
15	77035	1
16	77054	1
17	77065	1
18	77429	1
19	77479	1
20	77598	1
21	78017	1
22	78061	1
23	78336	1
24	785	1
25	78613	1
26	78734	1
27	78834	1
28	79520	1
29	79529	1
30	79553	1
31	79735	1
32	79761	1
33	79902	1

2.

```

texas_zip_closures = texas_zip_codes.merge(closures_by_zip, left_on='ZCTA5', right_on='ZIP Code', how='left')
#Fill nan with 0..
texas_zip_closures['Number of Closures'] = texas_zip_closures['Number of Closures'].fillna(0)

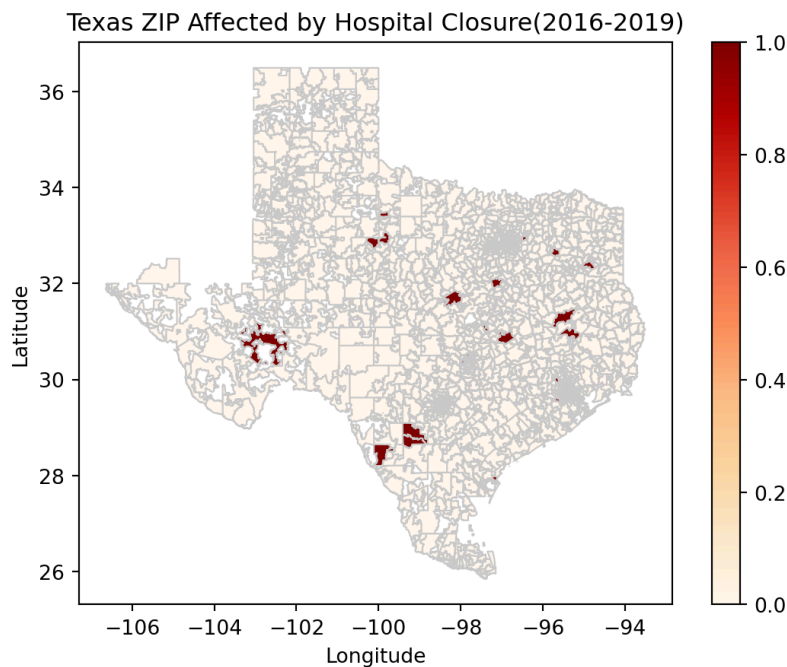
affected_zip_count = texas_zip_closures[texas_zip_closures['Number of Closures'] > 0]['ZCTA5'].nunique()
print(f"Directly affected ZIPs in Texas: {affected_zip_count}")

plt.figure(figsize=(12, 10))
texas_zip_closures.plot(column='Number of Closures', cmap='OrRd', linewidth=0.8, edgecolor='0.8', legend=True)
plt.title('Texas ZIP Affected by Hospital Closure(2016-2019)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()

```

Directly affected ZIPs in Texas: 33

<Figure size 1152x960 with 0 Axes>



3.

```
# 16093.44 meters
directly_affected_zips = texas_zip_closures[texas_zip_closures['Number of Closures'] > 0]
...

IMPORTANT!: print(directly_affected_zips.crs) to see the CRS! IT IS NOT METER.
...

directly_affected_zips = directly_affected_zips.to_crs("EPSG:3083")
texas_zip_codes = texas_zip_codes.to_crs("EPSG:3083")
# seems like this question doesn't need buffer around centroid.
...

directly_affected_zips['centroid'] = directly_affected_zips.geometry.centroid
directly_affected_zips['buffer'] = directly_affected_zips['centroid'].buffer(16093.4)
directly_affected_zips.set_geometry('buffer', inplace=True)
...

directly_affected_zips['geometry'] = directly_affected_zips.geometry.buffer(16093.4)
affected_zips = gpd.sjoin(texas_zip_codes, directly_affected_zips, how='inner', predicate='intersects')

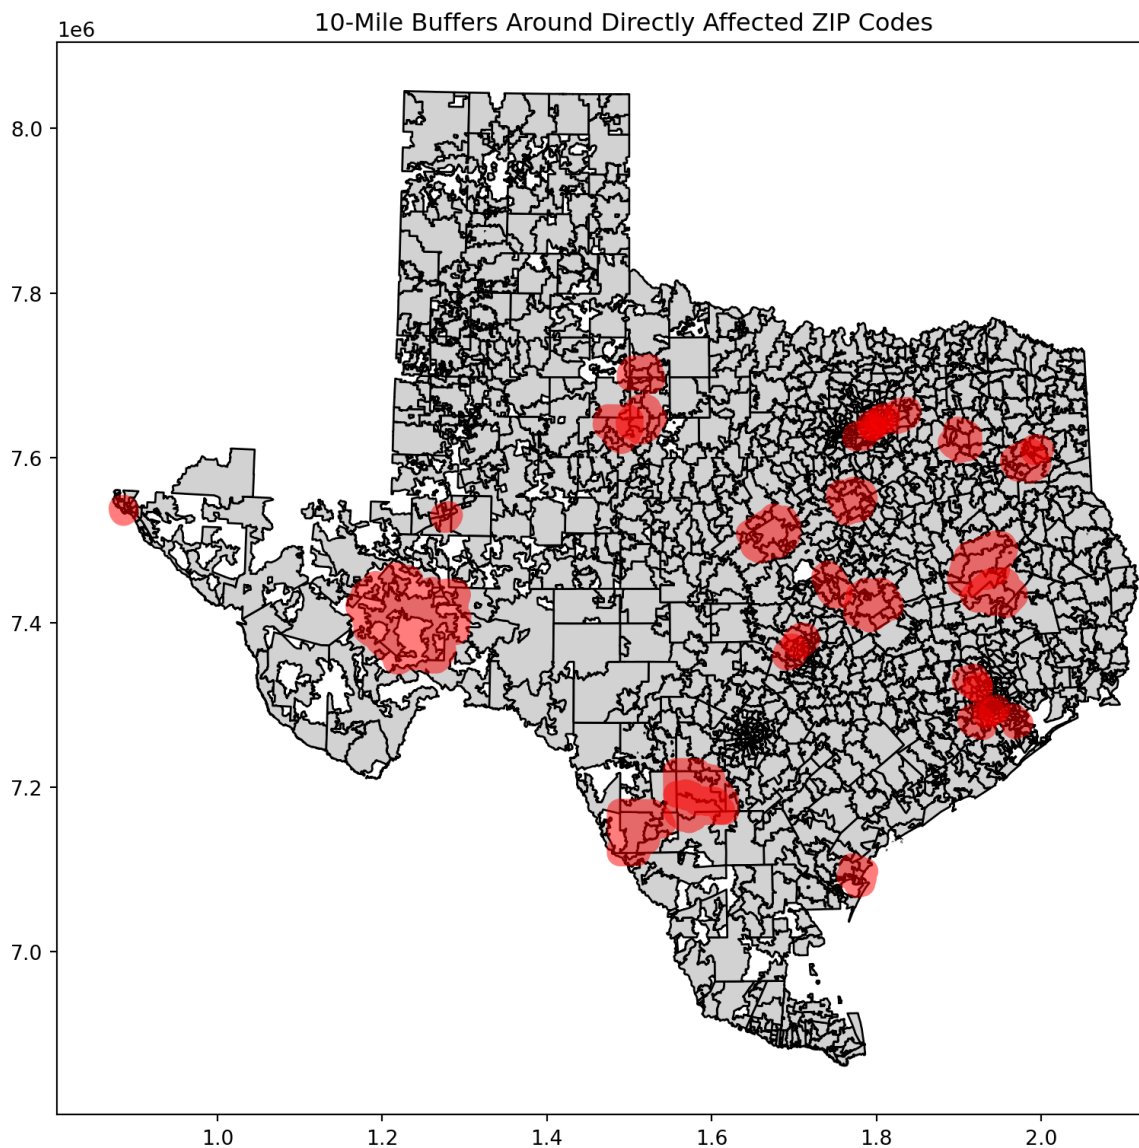
# sift out those directly; Must use ZCTA5_left, since the zcta5_right and ZIP Code is introduced from dir
indirectly_affected_zips = affected_zips[~affected_zips['ZCTA5_left'].isin(directly_affected_zips['ZCTA5_left'])]
indirect_zip_count = indirectly_affected_zips['ZCTA5_left'].nunique()
print("Number of indirectly affected ZIP codes:", indirect_zip_count)
```

Number of indirectly affected ZIP codes: 576

```
print(affected_zips.columns)
```

```
Index(['GEO_ID_left', 'ZCTA5_left', 'NAME_left', 'LSAD_left',
      'CENSUSAREA_left', 'geometry', 'index_right', 'GEO_ID_right',
      'ZCTA5_right', 'NAME_right', 'LSAD_right', 'CENSUSAREA_right',
      'ZIP Code', 'Number of Closures'],
      dtype='object')
```

```
# buffered areas
fig, ax = plt.subplots(figsize=(10, 10))
texas_zip_codes.plot(ax=ax, color='lightgrey', edgecolor='black')
directly_affected_zips.plot(ax=ax, color='red', alpha=0.5)
plt.title("10-Mile Buffers Around Directly Affected ZIP Codes")
plt.show()
```



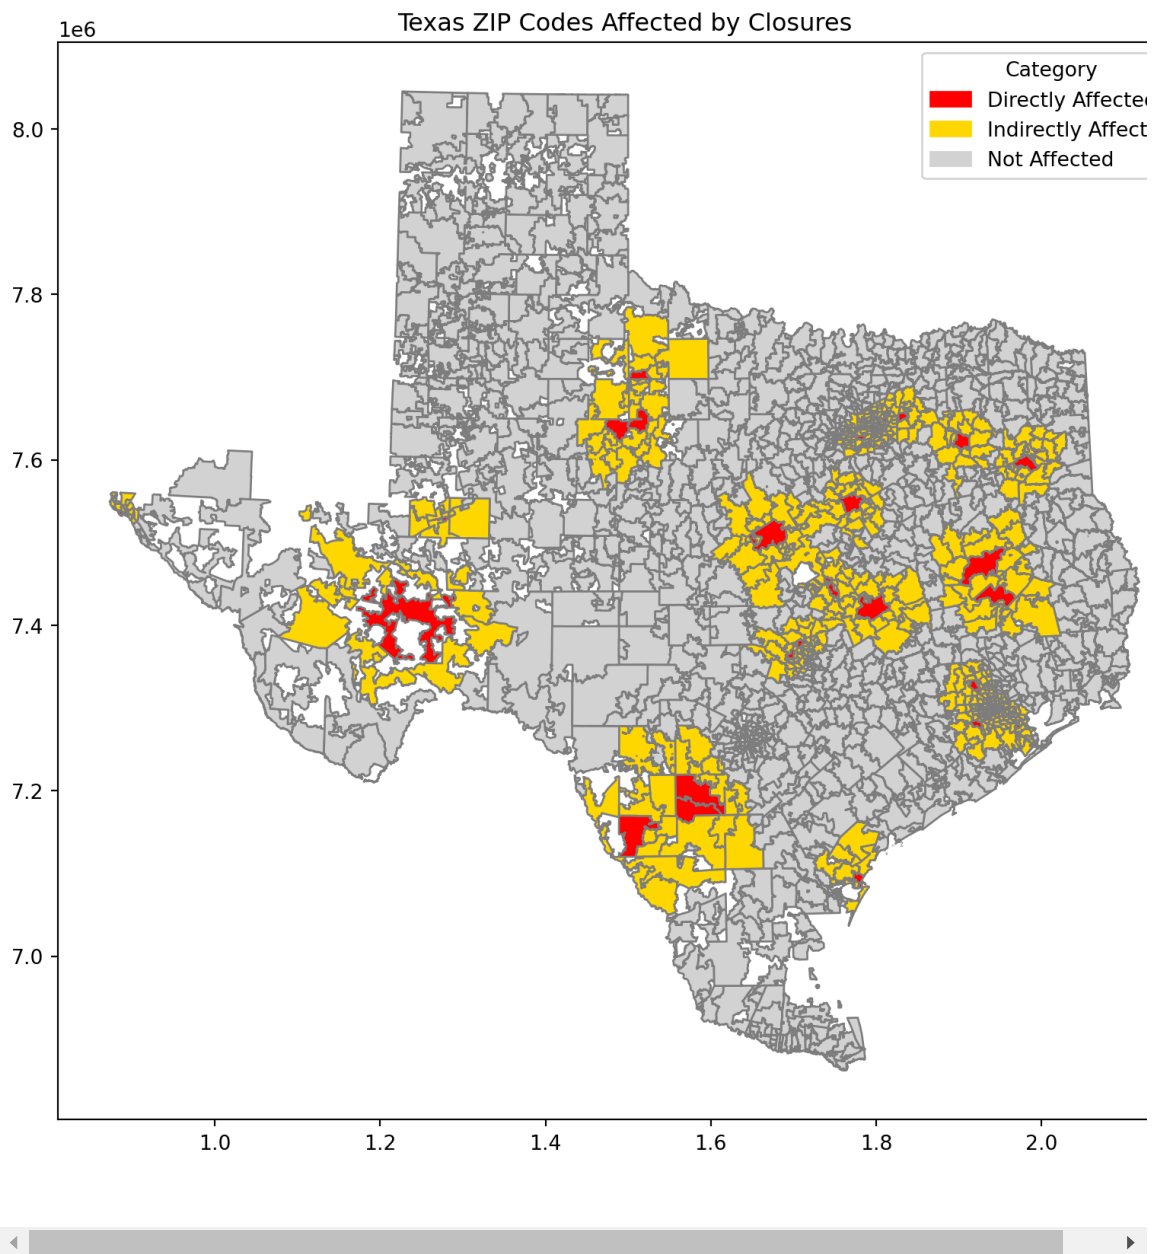
4.

```
def categorize_zip(zip):
    if zip['ZCTA5'] in directly_affected_zips['ZIP Code'].values:
        return 'Directly Affected'
    elif zip['ZCTA5'] in indirectly_affected_zips['ZCTA5_left'].values:
        return 'Indirectly Affected'
    else:
        return 'Not Affected'

# apply function onto dataframe, use .apply(func, axis=1):
texas_zip_codes['Category'] = texas_zip_codes.apply(categorize_zip, axis=1)
color_mapping = {
    'Directly Affected': 'red',
    'Indirectly Affected': 'gold',
    'Not Affected': 'lightgray'
}
texas_zip_codes['Color'] = texas_zip_codes['Category'].map(color_mapping)

fig, ax = plt.subplots(1, 1, figsize=(10, 10))
texas_zip_codes.plot(ax=ax, color=texas_zip_codes['Color'], edgecolor='gray', legend = True) #simply addi
ax.set_title('Texas ZIP Codes Affected by Closures')

#legend with custom patches
import matplotlib.patches as mpatches
legend_patches = [mpatches.Patch(color=color, label=label) for label, color in color_mapping.items()]
ax.legend(handles=legend_patches, title="Category")
plt.show()
```

Reflecting on the exercise (10 pts)

1.This is indeed doubtful. We are given the knowledge that when a TorF closure takes place, that will be reflected in the year's file. Thus, for an actual closure, the number of active hospital next year shall remain unaltered, and for false closure(merger/acquisition), the number next year should rise by 1. This indicates that the "non-decrease" criteria will fail to winnow out the false closures, they are all non-decrease! In a nutshell, the number we calculated in section2 is still the aggregate number of those 'simply non-active shown by termination code'. A remedy could be changing from 'not decrease' to 'increase', or sifting out those with more than 1 unique CMS number in the four years.

2.When we are identifying zip codes affected by closures,it may have duplicated date if the hospital is merged across different zip codes. For the hospital which located near the frontier of a zip code may happen that situation. I think the best way to improve this situation is by identifying the merge hospital with a 'merger_number',which may help us better identify the true closure.Instead of calculating distance alone, measure average travel time to the nearest hospital, considering public transportation, road quality, and traffic patterns. This would give a more realistic view of accessibility.Adjust the measure to account for population density and demographics. High-population or high-need areas could be weighted more heavily, reflecting a higher demand for healthcare services.